

## Assignment 3

In lab, we experimented with C++ code for initializing the PMU counters and retrieving the cyclecount. In this assignment, you'll be setting up your PMU counter to use in Python.

### Part A3.0: New kernel\_module code

- Download the new kernel\_module.zip code, go to canvas -> files -> assignments.
- Make and insert this new module following the ‘make’ and ‘insmod’ instructions from lab and in the README.
- Check that it is inserted on both CPUs by checking the dmesg | tail output
- I just followed the same instructions as the lab. It was very straightforward and I did not have any issues with building.

### Part A3.1: Access PMU from python

- Create a shared library object with two functions by wrapping the cycletime.h into a new shared opticed library (see Lab2)
  - One function to initialize the PMU counters
  - One function to get the cycle count
- Compile the shared library (see Lab2 if you’ve forgotten how to do this)
- Access the shared library functions using the ctypes module, **but don’t wrap the function calls in a python function.**
- I redefined the cycletime.h to be just a header file, and then moved the definitions of the functions into a cycletime.c file.
- To compile, I copied the Makefile over from the Lab and just ran make.
  - At first it didn’t work. I looked around online and edited the Makefile to make it work for the code and my own directory structure.
- I then created a jupyter notebook called test\_pmu. In there, I ran the two functions and made some simple code to get the cycle difference after sleeping for some arbitrary time.

### Part A3.2: Comparing and Gathering Data

In this section, we are going to use psutil to monitor CPU usage in percent, and the time module and PMU counting to evaluate the recursive fibonacci sequence timing operations.

- Isolate CPU 1 by editing the bootargs (see lab work part 1)
- Insert the CPUcntr kernel object onto both cpus using the instructions from lab
- Write code to do the following using the recur\_fibo function from lab
  - Initialize the cyclecounter
  - Get the ‘before’ time using the python time module
  - Get the ‘before’ cycle count
  - Run the recur\_fibo function on a CPU 1
  - Get the ‘after’ cycle count
  - Get the ‘after’ time count using the python time module
  - Get the cycle count and the amount of time used
- Vary the number of terms from 1 to 30 as you see fit to compare the different execution times
  - Take multiple trials for each variation (i.e. get three cyclecounts for n=5, then get three cyclecounts for n=10, etc) and average the different trials.
  - The error for each ‘n’ will be the standard deviation from the mean which is the standard deviation of all the trials divided by the square root of the number of trials.
- Plot the average results for varying ‘n’ along with error bars of your measurements.

- In order to compare the timing module and PMU counting, we need them to be in the same units.
  - To get the CPU frequency, run cat /proc/cpuinfo in a new terminal or run lscpu
  - Use this frequency to convert the PMU output from clock counts to timing
  - Compare the timing of the PMU counter to the timing module
- Before I started writing any code, I restructured my folder so everything stays organized.



```
xilinx@pynq:~/jupyter_notebooks/Assignment3$ ls -R
.:
kernel_module pmu_test python

./kernel_module:
CPUcntr.c CPUcntr.ko CPUcntr.mod CPUcntr.mod.c CPUcntr.mod.o CPUcntr.o Makefile
modules.order Module.symvers README

./pmu_test:
cycletime.c cycletime.h libcycletime.so Makefile

./python:
test_pmu.ipynb
```

- Using the recur\_fibo function, I added lines corresponding to each of the tasks in the bulleted list.
- I wrapped this in a function so I can recursively test the cycle and time counts with a different n every time.
- I then made a list of the n values I wanted to test
  - I chose n\_values = [5, 8, 10, 12, 15, 18, 20, 22, 25, 28, 30] for the following reasons:
    - Nothing less than 5 because the difference in computation is minimal and a little bit too fast for us to see anything
    - Equally spaced-ish intervals so we can see the exponentially longer computation time.
- I created some arrays to store the average time/cycle for each iteration of the test, and then ran a for loop that collected all of the data points and calculated the mean and error of each n value in both time and cycles.
  - I was originally considering a dictionary with a key value pair, where the key is the n value and the value is whatever I was trying to calculate. However that was unnecessary and I just used normal arrays for everything.
  - I changed the arrays into np arrays so calculations are easier and I can use the built in functions
- An issue that I hit while gathering the samples was that the cycle time would be 0. I checked though and it turned out it was because I wasn't isolating the CPU properly.
- For the graph, I used matplotlib, which has an error bar graph plotting function. I put in the corresponding values and plotted the results.



YouTube link: <https://youtu.be/torJvckuAA8>

Github: [https://github.com/cxu4426/WES237A-hw/tree/main/assignment\\_3](https://github.com/cxu4426/WES237A-hw/tree/main/assignment_3)

## Deliverables

Each student must submit the following individually

- A PDF report detailing your work flow and relative Jupyter notebook cells relating to your progress.
  - Your report should detail your work flow throughout the assignment. Be sure to discuss any difficulties or troubles you encountered and your troubleshooting procedure. Also, detail your thought process which led to the design and implementation of the code. For example, describe your top-down design methodology (i.e. how did you split the large task into smaller, more incremental jobs? How were you able to test each of these smaller parts?)
  - Please also include a plot with errorbars where the x-axis is the number of terms calculated (n) and the y-axis is the cyclecount
  - Please also include a plot with errorbars where the x-axis is the number of terms calculated (n) and the y-axis is the time taken to complete
- Your complete Jupyter notebook, downloaded as a PDF attached at the very end of your report
  - You can do this by selecting ‘File -> Print Preview’ then printing to PDF from the browser.
  - Use a PDF stitching tool like [pdfjoiner](#) to join your Report and Jupyter Notebook into a single PDF file

Each team should submit the following one per team

- All relevant code (.ipynb, .py, .cpp, .c, etc files) pushed to your team’s git repo.

```
1 #include <stdint.h>
2
3 /* Initialize PMU counters */
4 void pmu_init(int32_t do_reset, int32_t enable_divider)
5 {
6     int32_t value = 1;
7
8     if (do_reset)
9         value |= 6;      // reset all counters
10    if (enable_divider)
11        value |= 8;
12
13    value |= 16;
14
15    asm volatile ("MCR p15, 0, %0, c9, c12, 0\n\t" :: "r"(value));
16    asm volatile ("MCR p15, 0, %0, c9, c12, 1\n\t" :: "r"(0x8000000f));
17    asm volatile ("MCR p15, 0, %0, c9, c12, 3\n\t" :: "r"(0x8000000f));
18 }
19
20 /* Read cycle counter */
21 unsigned int pmu_get_cyclecount(void)
22 {
23     unsigned int value;
24     asm volatile ("MRC p15, 0, %0, c9, c13, 0\n\t" : "=r"(value));
25     return value;
26 }
27
```

```
In [1]: from ctypes import CDLL, c_uint, c_int
import time
```

```
In [2]: lib = CDLL("./libcycletime.so")

lib.pmu_init.argtypes = [c_int, c_int]
lib.pmu_get_cyclecount.restype = c_uint
```

```
In [19]: lib.pmu_init(1, 0)

start = lib.pmu_get_cyclecount()
time.sleep(0.5)
end = lib.pmu_get_cyclecount()

print("Cycle difference:", end - start)
```

Cycle difference: 9252565

## fib function

```
In [1]: def recur_fibo(n):
    if n <= 1:
        return n
    else:
        return(recur_fibo(n-1) + recur_fibo(n-2))
```

```
In [2]: from ctypes import CDLL, c_uint, c_int
import time
import matplotlib.pyplot as plt
import numpy as np
```

## implementation

```
In [3]: lib = CDLL("../pmu_test/libcycletime.so")

lib.pmu_init.argtypes = [c_int, c_int]
lib.pmu_get_cyclecount.restype = c_uint
```

```
In [4]: lib.pmu_init(1, 0)
```

```
Out[4]: 1
```

```
In [5]: def cycle_diff(c1, c0):
    MAX32 = 2**32
    if c1 >= c0:
        return c1 - c0
    else:
        return (MAX32 - c0) + c1
```

```
In [6]: def fibo_duration(n):
    t0 = time.time()
    c0 = lib.pmu_get_cyclecount()
    recur_fibo(n)
    c1 = lib.pmu_get_cyclecount()
    t1 = time.time()

    time_elapsed = t1 - t0
    cycle_elapsed = cycle_diff(c1, c0)
    return time_elapsed, cycle_elapsed
```

## gather cycles and times for different n for fibo function

```
In [7]: n_values = [5, 8, 10, 12, 15, 18, 20, 22, 25, 28, 30]
num_trials = 5
```

```
In [13]: time_avg = []
time_err = []
cycle_avg = []
cycle_err = []
```

```
for n in n_values:
    print(f"\n--- n = {n} ---")
    times = []
    cycles = []

    for trial in range(num_trials):
        t, c = fibo_duration(n)
        print(f"  trial {trial}: time = {t:.6f} s, cycles = {c}")
        times.append(t)
        cycles.append(c)

    times = np.array(times)
    cycles = np.array(cycles)

    # mean
    mean_time = times.mean()
    mean_cycle = cycles.mean()
    print(f"  mean_time = {mean_time:.6f} s")
    print(f"  mean_cycle = {mean_cycle}")

    time_avg.append(mean_time)
    cycle_avg.append(mean_cycle)

    # std
    std_time = np.sqrt(np.sum((times - mean_time) ** 2) / (num_trials - 1))
    time_err.append(std_time / np.sqrt(num_trials))

    std_cycle = np.sqrt(np.sum((cycles - mean_cycle) ** 2) / (num_trials - 1))
    cycle_err.append(std_cycle / np.sqrt(num_trials))
    print(f"  std_time = {std_time:.6f}, err_time = {time_err[-1]:.6f}")
    print(f"  std_cycle = {std_cycle:.6f}, err_cycle = {cycle_err[-1]:.6f}")
```

```
--- n = 5 ---
trial 0: time = 0.000074 s, cycles = 29718
trial 1: time = 0.000052 s, cycles = 25000
trial 2: time = 0.000051 s, cycles = 24479
trial 3: time = 0.000049 s, cycles = 23665
trial 4: time = 0.000048 s, cycles = 23708
mean_time = 0.000055 s
mean_cycle = 25314.0
std_time = 0.000011, err_time = 0.000005
std_cycle = 2524.355066, err_cycle = 1128.925905

--- n = 8 ---
trial 0: time = 0.000180 s, cycles = 88891
trial 1: time = 0.000143 s, cycles = 83242
trial 2: time = 0.000140 s, cycles = 82664
trial 3: time = 0.000140 s, cycles = 82484
trial 4: time = 0.000137 s, cycles = 80792
mean_time = 0.000148 s
mean_cycle = 83614.6
std_time = 0.000018, err_time = 0.000008
std_cycle = 3087.401626, err_cycle = 1380.727982

--- n = 10 ---
trial 0: time = 0.000339 s, cycles = 207238
trial 1: time = 0.001081 s, cycles = 687519
trial 2: time = 0.000329 s, cycles = 203582
trial 3: time = 0.000324 s, cycles = 201614
trial 4: time = 0.000324 s, cycles = 202431
mean_time = 0.000479 s
mean_cycle = 300476.8
std_time = 0.000336, err_time = 0.000150
std_cycle = 216373.850543, err_cycle = 96765.327673

--- n = 12 ---
trial 0: time = 0.000823 s, cycles = 522679
trial 1: time = 0.000807 s, cycles = 516022
trial 2: time = 0.000805 s, cycles = 514414
trial 3: time = 0.000812 s, cycles = 519169
trial 4: time = 0.000806 s, cycles = 515516
mean_time = 0.000811 s
mean_cycle = 517560.0
std_time = 0.000007, err_time = 0.000003
std_cycle = 3362.801585, err_cycle = 1503.890588

--- n = 15 ---
trial 0: time = 0.003504 s, cycles = 2264527
trial 1: time = 0.003372 s, cycles = 2182580
trial 2: time = 0.003571 s, cycles = 2309944
trial 3: time = 0.003363 s, cycles = 2176923
trial 4: time = 0.003356 s, cycles = 2173125
mean_time = 0.003433 s
mean_cycle = 2221419.8
std_time = 0.000098, err_time = 0.000044
std_cycle = 62280.911319, err_cycle = 27852.870282

--- n = 18 ---
trial 0: time = 0.014334 s, cycles = 9303733
trial 1: time = 0.014102 s, cycles = 9155820
trial 2: time = 0.014169 s, cycles = 9199638
trial 3: time = 0.014101 s, cycles = 9155804
```

```
trial 4: time = 0.014187 s, cycles = 9211677
mean_time = 0.014179 s
mean_cycle = 9205334.4
std_time = 0.000095, err_time = 0.000043
std_cycle = 60539.033609, err_cycle = 27073.878888

--- n = 20 ---
trial 0: time = 0.036955 s, cycles = 24007753
trial 1: time = 0.039717 s, cycles = 25801237
trial 2: time = 0.053369 s, cycles = 34671425
trial 3: time = 0.042405 s, cycles = 27547825
trial 4: time = 0.036910 s, cycles = 23979806
mean_time = 0.041871 s
mean_cycle = 27201609.2
std_time = 0.006817, err_time = 0.003049
std_cycle = 4428876.523748, err_cycle = 1980653.794211

--- n = 22 ---
trial 0: time = 0.096583 s, cycles = 62764395
trial 1: time = 0.114845 s, cycles = 74630834
trial 2: time = 0.098347 s, cycles = 63907237
trial 3: time = 0.096155 s, cycles = 62489630
trial 4: time = 0.115354 s, cycles = 74963700
mean_time = 0.104257 s
mean_cycle = 67751159.2
std_time = 0.009934, err_time = 0.004442
std_cycle = 6455190.057890, err_cycle = 2886848.755425

--- n = 25 ---
trial 0: time = 0.427066 s, cycles = 277572020
trial 1: time = 0.433800 s, cycles = 281949555
trial 2: time = 0.511557 s, cycles = 332492205
trial 3: time = 0.568255 s, cycles = 369340791
trial 4: time = 0.426703 s, cycles = 277338097
mean_time = 0.473476 s
mean_cycle = 307738533.6
std_time = 0.063931, err_time = 0.028591
std_cycle = 41553756.960219, err_cycle = 18583405.056711

--- n = 28 ---
trial 0: time = 1.749498 s, cycles = 1137152977
trial 1: time = 1.753391 s, cycles = 1139682760
trial 2: time = 1.749792 s, cycles = 1137345951
trial 3: time = 1.746871 s, cycles = 1135446388
trial 4: time = 1.761321 s, cycles = 1144837000
mean_time = 1.752174 s
mean_cycle = 1138893015.2
std_time = 0.005615, err_time = 0.002511
std_cycle = 3648802.261836, err_cycle = 1631793.978784

--- n = 30 ---
trial 0: time = 4.782565 s, cycles = 3108644274
trial 1: time = 4.549037 s, cycles = 2956855422
trial 2: time = 4.606897 s, cycles = 2994460727
trial 3: time = 4.790235 s, cycles = 3113632050
trial 4: time = 4.556496 s, cycles = 2961702639
mean_time = 4.657046 s
mean_cycle = 3027059022.4
std_time = 0.120193, err_time = 0.053752
std_cycle = 78124409.773294, err_cycle = 34938298.191027
```

plot!

```
In [14]: freq = 650e6
```


```
cycle_time_avg = np.array(cycle_avg) / freq
cycle_time_err = np.array(cycle_err) / freq
```


```
In [17]: import matplotlib.pyplot as plt
```

```
plt.figure()
plt.errorbar(n_values, cycle_time_avg, yerr=cycle_time_err, fmt='o-', capsizes=5)
plt.yscale("log")
plt.xlabel("Terms Calculated (n)")
plt.ylabel("Time (seconds)")
plt.title("Cyclecount Converted to Seconds VS Terms Calculated")
plt.grid(True)
plt.show()

plt.figure()
plt.errorbar(n_values, cycle_avg, yerr=cycle_err, fmt='o-', capsizes=5)
plt.yscale("log")
plt.xlabel("Terms Calculated (n)")
plt.ylabel("Cyclecount")
plt.title("Cyclecount VS Terms Calculated")
plt.grid(True)
plt.show()

plt.figure()
plt.errorbar(n_values, time_avg, yerr=time_err, fmt='o-', capsizes=5)
plt.yscale("log")
plt.xlabel("Terms Calculated (n)")
plt.ylabel("Time (seconds)")
plt.title("Time Vs Terms Calculated")
plt.grid(True)
plt.show()
```





In [ ]: